

Contents

Bayesian Filters for Dynamic State Space Models	2
Inference, Estimation and Forecasting	3
Importance and Sequential Importance Sampling	4
Example 1: Simple Weighted Points Examples	6
Sequential Importance Sampling (SIS) for State Space Models	6
Example 2: SIS Filter for a Random Walk Measured With Noise	10
Final Thoughts For Now	11
Bayesian Filters for Static Parameters in Dynamic Models	13

Bayesian Filters for Dynamic State Space Models

Consider a discrete time system of states \mathbf{x}_t which evolves according to a dynamic model

$$\mathbf{x}_t = m(\mathbf{x}_{t-1}, \epsilon_t) \quad (1)$$

where ϵ_t is a random variable representing uncertainty associated with \mathbf{x}_t at time t . These states are not directly observed, but instead we have a sequence of observations $\{y_1, y_2, \dots, y_t\}$ which were generating **using** the states according to

$$\mathbf{y}_t = h(\mathbf{x}_t, \eta_t) \quad (2)$$

where η_t is a random variable describing measurement uncertainty (or some other type of uncertainty contained in the observation sequence). Both \mathbf{x}_t and \mathbf{y}_t are vector variables in general but the boldface will be omitted for clarity until it is needed. The sequences $\{x_1, x_2, \dots, x_t\}$ and $\{y_1, y_2, \dots, y_t\}$ are ordered (ie y_2 was observed *before* y_3 and are referred to using the following index notation $x_{1:t} = \{x_1, x_2, \dots, x_t\}$ and similar for y . Finally, the state sequence is appended with an initial condition x_0 and $p(x_0)$ which is assumed to be known.

The fundamental problem in the context of estimation and forecasting is to calculate at time t :

- $p(x_t | y_{1:t-1})$ ie our belief about the current state x_t **prior** to receiving a contemporaneous observation y_t
- $p(x_t | y_{1:t})$ ie our belief about the current state x_t **after** receiving a current observation y_t

Iterative Monte Carlo technique are required for solving the integrals that arise when computing these quantities in the Bayesian context. In the case where the transitions are linear and Gaussian the recursions derived can be solved analytically - these are the Kalman Filter solutions. More generally specialised numerical techniques are required.

The following equations are deliberately recursive so that they can be updated online. To derive them the *probability chain rule* (or telescope rule) is used extensively ie

$$\begin{aligned} p(x_1, x_2, \dots, x_n) &= p(x_n | x_{n-1}, \dots, x_1) p(x_{n-1}, \dots, x_0) \\ &= p(x_n | x_{n-1}, \dots, x_1) p(x_{n-1} | x_{n-2}, \dots, x_0) p(x_{n-2}, \dots, x_1) \\ &= p(x_n | x_{n-1}, \dots, x_1) p(x_{n-1} | x_{n-2}, \dots, x_0) \dots p(x_2 | x_1) p(x_1) \end{aligned}$$

In the context of dynamic models the subscripts above will be a time index and the data $\{x_1, \dots, x_t\}$ are ordered (ie x_{t-1} occurs before x_t). Although the expression *transition kernel* has specific meaning when applied to dynamic systems I use the expression to refer to any expression which modifies an expression available at time $t - 1$ to one at t . In the above I would refer to

$$\begin{aligned} p(x_{1:t}) &= p(x_t | x_{1:t-1}) \dots p(x_2 | x_1) p(x_1) \\ &= k_{t-1,t} p(x_t | x_{1:t-1}) \end{aligned}$$

where

$$k_{t-1,t} = p(x_{t-1} | x_{1:t-2}) \dots p(x_2 | x_1) p(x_1) \quad (3)$$

is the *transition kernel* which brings the expression $p(x_{t-1} | x_{1:t-2})$ to one at time t . As we are developing online filtering methods the calculation of transition kernels is a key focus of these notes.

Inference, Estimation and Forecasting

The two distributions $p(x_t|y_{1:t-1})$ and $p(x_t|y_{1:t})$ are computed using a system of non-linear equations that are quickly derived from the state equations

$$\begin{aligned}x_t &= m(x_{t-1}, \epsilon_t) \\ y_t &= h(x_t, \eta_t)\end{aligned}$$

and a couple of probability identities. First consider $p(x_t|y_{1:t-1})$ which represents our **prior** for x_t before receiving y_t :

$$\begin{aligned}p(x_t|y_{1:t-1}) &= \int p(x_t|x_{t-1}, y_{1:t-1})p(x_{t-1}|y_{1:t-1})dx_{t-1} \\ &= \int p(x_t|x_{t-1})p(x_{t-1}|y_{1:t-1})dx_{t-1}\end{aligned}$$

where in the last line we used $p(x_t|x_{t-1}, y_{1:t-1}) = p(x_t|x_{t-1})$, noting that this does *not* mean that x_t is independent of the history $y_{1:t-1}$, but rather that the history $y_{1:t-1}$ is redundant when the prior state x_{t-1} is known. Also, we marginalise over x_{t-1} in the first line, but note that we need to condition on the history $y_{1:t-1}$ when integrating out x_{t-1} , ie the $p(x_{t-1}|y_{1:t-1})$ term.

After receiving y_t we want to update and calculate $p(x_t|y_{1:t})$ which is done as follows:

$$\begin{aligned}p(x_t|y_{1:t}) &= \frac{p(x_t, y_{1:t})}{p(y_{1:t})} \\ &= \frac{p(y_t|x_t, y_{1:t-1})p(x_t, y_{1:t-1})}{p(y_t|y_{1:t-1})p(y_{1:t-1})} \\ &= \frac{p(y_t|x_t, y_{1:t-1})p(x_t|y_{1:t-1})p(y_{1:t-1})}{p(y_t|y_{1:t-1})p(y_{1:t-1})} \\ &= \frac{p(y_t|x_t, y_{1:t-1})p(x_t|y_{1:t-1})}{p(y_t|y_{1:t-1})}\end{aligned}$$

In the denominator marginalise over x_{t-1} in a similar way to what we did in the prior distribution ie

$$\begin{aligned}p(y_t|y_{1:t-1}) &= \int p(y_t|x_t, y_{1:t-1})p(x_t|y_{1:t-1})dx_t \\ &= \int p(y_t|x_t)p(x_t|y_{1:t-1})dx_t\end{aligned}$$

where we used $p(y_t|x_t, y_{1:t-1}) = p(y_t|x_t)$ again. Combining this result in the denominator with our previous prior we have the required system of equations:

$$\begin{aligned}p(x_t|y_{1:t-1}) &= \int p(x_t|x_{t-1})p(x_{t-1}|y_{1:t-1})dx_{t-1} \\ p(x_t|y_{1:t}) &= \frac{p(y_t|x_t, y_{1:t-1})p(x_t|y_{1:t-1})}{\int p(y_t|x_t)p(x_t|y_{1:t-1})dx_t}\end{aligned}$$

To demonstrate consider the online inference problem where we first receive y_1 and then y_2 . Our algorithm is as follows.

- Step 1

Prior to any observations $y = \{\}$ is an empty set and we are relying on our initial condition $p_0(x_0)$

which is assumed known. Our prior is

$$\begin{aligned} p(x_1 | \{\}) &= \int p(x_1 | x_0) p(x_0 | \{\}) dx_0 \\ &= \int p(x_1 | x_0) p_0(x_0) dx_0 \end{aligned}$$

with $p_0(x_0)$ the assumed initial condition and $p(x_1 | x_0)$ is determined by the known system update equations.

- Step 2

On arrival of y_1 we update the above prior to calculate $p(x_1 | y_1)$ as follows:

$$p(x_1 | y_1) = \frac{p(y_1 | x_1) p(x_1 | \{\})}{\int p(y_1 | x_1) p(x_1 | \{\}) dx_1}$$

where $p(y_1 | x_1)$ is determined by the known system update equations and $p(x_1 | \{\})$ is the prior we calculated in Step 1 above.

- Step 3

After receiving y_1 and prior to receiving y_2 we can again compute a prior

$$p(x_2 | y_{1:1}) = \int p(x_2 | x_1) p(x_1 | y_{1:1}) dx_1 \quad (4)$$

where $p(x_2 | x_1)$ is a system update equation and $p(x_1 | y_{1:1})$ was calculated in step 2.

- Step 4

Finally we receive y_2 , which we use to update our above prior:

$$p(x_2 | y_{1:2}) = \frac{p(y_{1:2} | x_2) p(x_2 | y_{1:1})}{\int p(y_{1:2} | x_2) p(x_2 | y_{1:1}) dx_2} \quad (5)$$

where $p(x_2 | y_{1:1})$ was the prior from Step 3 above and $p(y_{1:2} | x_2)$ is from the system update equations.

This sequence continues. As a final illustration of the point assume that the dynamics equations where

$$\begin{aligned} y_t &= 2x_t + \epsilon_t \\ x_t &= x_{t-1} + \eta_t \end{aligned}$$

where ϵ_t and η_t where $\mathcal{N}(0, \sigma_\epsilon^2)$ and $\mathcal{N}(0, \sigma_\eta^2)$ so that we have a linear, Gaussian system. In Step 3 for instance we need $p(x_2 | x_1)$, which we can calculate from the second equation above as $p(x_2 | x_1) = \mathcal{N}(x_1, \sigma_\eta^2)$, and $p(y_2 | x_2)$ in Step 4 would be $p(y_2 | x_2) = \mathcal{N}(2x_2, \sigma_\epsilon^2)$.

As mentioned earlier in special cases the above system can be solved analytically, an example being the Kalman solution for linear Gaussian systems. The remainder of these notes will investigate how to approximate these distributions for general systems.

Importance and Sequential Importance Sampling

Sequential Importance Sampling (SIS) is the first step on the way to computing the difficult integrals arising from the above expressions. It is an iterative technique for computing the joint distribution $p(x_{0:t} | y_{1:t})$. To get there, first consider the standard Importance Sampling Problem.

Given a function $x \in \mathbb{R}$ define the expectation

$$\mathbb{E}_p(f(x)) = \int f(x)p(x)dx$$

where \mathbb{E}_p denotes expectation with respect to the distribution p . The above integral may be difficult or impossible to compute when the above integral is not available in closed form (or is high dimensional...). Define the estimate

$$\hat{I} = \frac{1}{N} \sum f(x_i) \quad x_i \sim p(x)$$

where x_i are all *iid* draws from the distribution $p(x)$ and notice that the estimator is unbiased

$$\begin{aligned} \mathbb{E}_p(\hat{I}) &= \mathbb{E}_p(\sum f(x_i)/N) \\ &= \frac{1}{N} \cdot N \cdot \mathbb{E}_p(f(x)) \\ &= \mathbb{E}_p(f(x)) \end{aligned}$$

The problem arises in the situation that

- we cannot sample $x \sim p(x)$, either because we do not know how or the process required to generate samples is prohibitively expensive. Or
- The approximation may exhibit high variance if $x \sim p(x)$ generates a large fraction of samples that do not contribute greatly to the integral (ie provide much information about the function).

Instead assume that despite not knowing how to sample from $p(x)$ we can still evaluate it, and that we have an importance distribution $g(x)$ with the same domain as $p(x)$ that is easy to sample from. Perform the following trick

$$\begin{aligned} \mathbb{E}_p(f(X)) &= \int f(x)p(x)dx \\ &= \int \frac{f(x)p(x)}{g(x)}g(x)dx \\ &= \int r(x)g(x)dx \\ &= \mathbb{E}_g(r(x)) \end{aligned}$$

to see we can either consider our integral to be a sample average of $f(x)$, where x is drawn from $p(x)$ (ie estimate the expectation \mathbb{E}_p) or a sample average of the new function $r(x) = f(x)p(x)/g(x)$ where x is drawn from the importance distribution $g(x)$ (ie estimate the expectation \mathbb{E}_g). There is much to be said about how to choose $g(x)$ which I will gloss over here. Robert & Casella (2004) have a good description and further reading on this topic.

In our context (state space models) the distribution we are required to sample from (ie $p(x)$ in this case) is changing at each iteration, which means we are constantly having to tune our importance distribution $g(x)$. Consistent with the iterative ideas introduced earlier we look for a sequential way to build up these distributions - **Sequential Importance Sampling**. A final point to be made is that SIS is focused on the notation of *weights* in the above MC integrals, which I guess is a third

way to interpret the expectation above ie

$$\begin{aligned}
\mathbb{E}_p(f(x)) &= \int f(x)p(x)dx \\
&= \int \frac{p(x)}{g(x)}f(x)g(x)dx \\
&= \int w(x)f(x)g(x)dx \\
&= \mathbb{E}_g(w(x)f(x))
\end{aligned}$$

where $w(x)$ is a weight function representing a reweighting of the integral because we are taking an expectation with $g(x)$ instead of $p(x)$ as the problem requires. The estimator for this integral would be

$$\hat{I} = \sum \hat{w}(x_i)f(x_i) \quad x_i \sim g(x) \quad (6)$$

where $\hat{w}(x)$ can be either

$$\hat{w}(x) = \frac{w(x_i)}{N} \quad (7)$$

or

$$\hat{w}(x) = \frac{w(x_i)}{\sum w(x_i)} \quad (8)$$

although the second form (usually referred to as *self normalising weights*) is more common, because the former requires the importance distribution $g(x)$ to be known completely while the latter allows it to be known only up to a multiplicative constant. In either case we are *drawing* samples $\{x_i, w(x_i)\}$ with which we calculate our weighted averages for estimating integrals. The Particle Filter literature will refer to this sample of point-weight pairs $\{x_i, w(x_i)\}$ as **particles** and their respective **weights**.

Example 1: Simple Weighted Points Examples

Here is a simple example where we use a weighted set of particles generated from one distribution to approximate an other. Here my importance distribution q is the two dimensional uniform distribution on $(-4, 4)$, and my target p is the bivariate normal with mean zero, variance one and correlation 0.5. I am not considering separate state and observation equations so there is no need for conditioning in these distributions. I compute weights w_i for each point as

$$\begin{aligned}
w_i &= \frac{p(x_i)}{q(x_i)} \\
&= p(x_i)
\end{aligned}$$

leaving out the uniform normalisation constant $1/64$ as I use normalised weights anyway. The following plot has the uniform importance distribution in red and the target $p(x_i)$ in blue. The points are scaled in size based on their weights.

Sequential Importance Sampling (SIS) for State Space Models

In the state space context we are given a function $f(x_{0:t})$ of our states, and are computing the expectation

$$\mathbb{E}(f(x_{0:t})) = \int f(x_{0:t})p(x_{0:t}|y_{1:t})dx_{0:t} \quad (9)$$

This integral using MC techniques is desirable because it converts what is likely to be a difficult integration problem into a summation problem. This is difficult to compute however since $p(x_{0:t}|y_{1:t})$ cannot in general be sampled from. Consistent with the IS approach previously we




Figure 1: Approximating a Target Distribution with Weighted Points

instead calculate

$$\mathbb{E}(f(x_{0:t})) = \int f(x_{0:t})w(x_{0:t}|y_{1:t})g(x_{0:t}|y_{1:t})dx_{0:t} \quad (10)$$

where the *importance weights* are

$$w(x_{0:t}|y_{1:t}) = \frac{p(x_{0:t}|y_{1:t})}{g(x_{0:t}|y_{1:t})} \quad (11)$$

In practice we draw a set of N point-weight pairs $\{x_{0:t}^i, w(x_{0:t}^i|y_{1:t})\}$ where $i = 1, 2, \dots, N$ and estimate the integral as expected

$$\mathbb{E}(f(x_{0:t})) = \sum \hat{w}^i(x_{0:t}|y_{1:t})f(x_{0:t}^i) \quad \hat{w}^i(x_{0:t}|y_{1:t}) = \frac{w^i(x_{0:t}|y_{1:t})}{\sum w^i(x_{0:t}|y_{1:t})} \quad (12)$$

Nothing here is conceptually different to the previous more simple case except we know have lot more notation. The task at hand now is to build the (posterior) distribution $p(x_{0:t}|y_{1:t})$ using a MC approach, since calculus based approaches are almost certainly going to be infeasible.

The SIS approach, of which particle filtering is a particular example, is to build up an approximation to the posterior in an iterative way using the set of N point-weight pairs $\{x_{0:t}^i, w(x_{0:t}^i|y_{1:t})\}$. Effectively, we are at each point in time approximating the posterior as a summation:

$$p(x_{0:t}|y_{1:t}) = w_{0:t}^i \delta(x_{0:t} - x_{0:t}^i) \quad (13)$$

where the weights w are carefully constructed in an iterative way discussed below and the delta function $\delta(x_{0:t} - x_{0:t}^i)$ represents a point mass at each point in our sample of point-weight pairs. At each iteration we have new points and new weights, and these points create a *trajectory* of N points over time

$$\{x_0\}_1^N \rightarrow \{x_{0:1}\}_1^N \rightarrow \dots \rightarrow \{x_{0:t}\}_1^N \quad (14)$$

Again, the notation in the weight for example $w_{0:t}^i$ signifies that for sample i the weight function at time t is built from data all the way from $t = 0$ to t , and similar for the x above.

The challenge remains to work out:

- how to move our samples between iterations (ie generate x_t using information from $t - 1$)
- how to update weights w_t^i using information from $t - 1$.

To proceed write the posterior $p(x_{0:t}|y_{1:t}) = p(x_{0:t}, y_{1:t}) / p(y_{1:t})$ and expand using the chain rule repeatedly to expose transition distributions like $p(y_t|x_t)$ and $p(x_t|x_{t-1})$ as follows:

$$\begin{aligned}
p(x_{0:t}|y_{1:t}) &= \frac{p(x_{0:t}, y_{1:t})}{p(y_{1:t})} \\
&= \frac{p(y_t|y_{1:t-1}, x_{0:t})p(x_t|y_{1:t-1}, x_{0:t-1})p(y_{1:t-1}, x_{0:t-1})}{p(y_t|y_{1:t-1})p(y_{1:t-1})} \\
&= \frac{p(y_t|y_{1:t-1}, x_{0:t})p(x_t|y_{1:t-1}, x_{0:t-1})p(x_{0:t-1}|y_{1:t-1})p(y_{1:t-1})}{p(y_t|y_{1:t-1})p(y_{1:t-1})} \\
&= \frac{p(y_t|y_{1:t-1}, x_{0:t})p(x_t|y_{1:t-1}, x_{0:t-1})p(x_{0:t-1}|y_{1:t-1})}{p(y_t|y_{1:t-1})} \\
&\rightarrow p(x_{0:t}|y_{1:t}) = K_{t-1,t}p(x_{0:t-1}|y_{1:t-1})
\end{aligned}$$

where the kernel $K_{t-1,t}$ is defined

$$K_{t-1,t} = \frac{p(y_t|y_{1:t-1}, x_{0:t})p(x_t|y_{1:t-1}, x_{0:t-1})}{p(y_t|y_{1:t-1})} \quad (15)$$

and is the way the posterior updates from $t - 1$ to t . The recursive nature of the problem is clear. Also, the notation can be simplified further using

$$\begin{aligned}
p(y_t|y_{1:t-1}, x_{0:t}) &= p(y_t|x_t) \\
p(x_t|y_{1:t-1}, x_{0:t-1}) &= p(x_t|x_{t-1})
\end{aligned}$$

to write

$$K_{t-1,t} = \frac{p(y_t|x_t)p(x_t|x_{t-1})}{p(y_t|y_{t-1})} \quad (16)$$

Combine the usual expression for the weight function with our iterative importance function above to first write

$$\begin{aligned}
w(x_{0:t}|y_{1:t}) &= \frac{p(x_{0:t}|y_{1:t})}{g(x_{0:t}|y_{1:t})} \\
&= \frac{p(x_{0:t}|y_{1:t})}{g(x_t|x_{0:t-1}, y_{1:t})g(x_{t-1}|y_{1:t-1})} \\
&= \frac{p(x_{0:t}|y_{1:t})}{g(x_t|x_{t-1}, y_t)g(x_{t-1}|y_{1:t-1})}
\end{aligned}$$

where the denominator is equation (10) in Creale (2009), but I don't understand how the second expression is condition on $y_{1:t-1}$ and not $y_{1:t}$. Substitute for $p(x_{0:t}|y_{1:t})$ using the expression with

the kernel $K_{t-1,t}$ above and simplify

$$\begin{aligned}
w(x_{0:t}|y_{1:t}) &= \frac{p(x_{0:t}|y_{1:t})}{g(x_t|x_{t-1}, y_t)g(x_{0:t-1}|y_{1:t-1})} \\
&= \frac{K_{t-1,t}p(x_{0:t-1}|y_{1:t-1})}{g(x_t|x_{t-1}, y_t)g(x_{t-1}|y_{1:t-1})} \\
&= \frac{p(y_t|x_t)p(x_t|x_{t-1})}{p(y_t|y_{t-1})} \frac{p(x_{t-1}|y_{t-1})}{g(x_t|x_{t-1}, y_t)g(x_{t-1}|y_{1:t-1})} \\
&= \frac{p(y_t|x_t)p(x_t|x_{t-1})}{g(x_t|x_{t-1}, y_t)p(y_t|y_{t-1})} \frac{p(x_{t-1}|y_{t-1})}{g(x_{t-1}|y_{t-1})} \\
&\propto \frac{p(y_t|x_t)p(x_t|x_{t-1})}{g(x_t|x_{t-1}, y_t)} \frac{p(x_{t-1}|y_{t-1})}{g(x_{t-1}|y_{t-1})} \\
\rightarrow w(x_{0:t}|y_{1:t}) &= \hat{w}_{t-1,t} \cdot w_{t-1}
\end{aligned}$$

As before $p(y_t|x_t)$ and $p(x_t|x_{t-1})$ come from the known system equations. In the denominator the second last equation the $p(y_t|y_{t-1})$ term was dropped and the expression reduced to one of proportionality. This can be done as we use self-normalising weights but I'm not sure the exact reason for this but I presume the reasoning is the same as a usual Bayesian statement that $p(A|B) \propto p(B|A)p(A)$ where the denominator is hard to calculate and in general is not needed.

How do we use this all in practice. Remember we are trying to construct a sequence:

$$\{x_0\}_1^N \rightarrow \{x_{0:1}\}_1^N \rightarrow \dots \rightarrow \{x_{0:t}\}_1^N \quad (17)$$

where at each time point are point-weight samples $\{x_{0:t}, w_{0:t}\}$ are sampled from our importance distribution g using the sequential form

$$g(x_{0:t}|y_{1:t}) = g(x_t|x_{0:t-1}, y_{1:t})g(x_{0:t-1}|y_{1:t-1}) \quad (18)$$

The second term $g(x_{0:t-1}|y_{1:t-1})$ is a sum of delta functions centred at each point in our sample at time $t-1$ ie $\{x_{0:t-1}\}_1^N$, and the first term $g(x_t|x_{0:t-1}, y_{1:t})$ is the distribution we draw from to construct our new sample $\{x_{0:t}\}_1^N$. With these new points we recompute their respective weights using

$$w(x_{0:t}|y_{1:t}) = \hat{w}_{t-1,t} \cdot w_{t-1} \quad (19)$$

where points where with large $\hat{w}_{t-1,t}$ ie

$$\hat{w}_{t-1,t} = \frac{p(y_t|x_t)p(x_t|x_{t-1})}{g(x_t|x_{t-1}, y_t)} \quad (20)$$

occurs when $p(y_t|x_t)p(x_t|x_{t-1})$ are large. This is as expected - we generated the x_t from a importance distribution g but we should assign large weights to samples where these transitions $p(y_t|x_t)$ and $p(x_t|x_{t-1})$, both of which are determined by the system dynamics, were most likely. At this point we are in a similar position to standard IS:

$$\mathbb{E}(f(x_{0:t})) = \sum \hat{w}_t^i f(x_{0:t}^i) \quad \hat{w}_t^i = \frac{w_t^i}{\sum w_t^i} \quad (21)$$

and the posterior itself is estimated as

$$p(x_{0:t}|y_{1:t}) = \sum \hat{w}_t^i \delta(x_{0:t} - x_{0:t}^i) \quad (22)$$

From here we have not spoken about two main points:

- What determines good choices of importance distribution q ?
- Do we run into issues with this process after many iterations?

Example 2: SIS Filter for a Random Walk Measured With Noise

Consider the following simple state space model:

$$\begin{aligned}x_t &= \mu + x_{t-1} + \eta_t & \eta_t &\sim \mathbb{N}(0, \Sigma_\eta) \\ y_t &= x_t + \epsilon_t & \epsilon_t &\sim \mathbb{N}(0, \Sigma_\epsilon)\end{aligned}$$

where the state $x_t = [x_{1,t}, x_{2,t}]$ is a random walk in two dimensions with drift. It is measured $y_t \in \mathbb{R}^2$ with noise represented by Σ_ϵ . The code for this example is in `scripts/ex2.py`. A sample trajectory is plotted below. Remember in a practical application the state path x_t (blue points) is not observed, instead only the observations y_t are known (red points).




Figure 2: Example Path for Example 2 SIS Filter

To apply the SIS filter choose importance density $g(x_t|x_{0:t-1}, y_{1:t}) = p(x_t|x_{t-1})$, which from the state equations is simple

$$p(x_t|x_{t-1}) \sim \mathbb{N}(\mu + x_{t-1}, \eta_t) \quad (23)$$

and this choice simplifies the weight calculation dramatically

$$\begin{aligned}w_t &= \hat{w}_{t-1,t} \cdot w_{t-1} \\ &= \frac{p(y_t|x_t)p(x_t|x_{t-1})}{g(x_t|x_{t-1})} \cdot w_{t-1} \\ &= p(y_t|x_t) \cdot w_{t-1}\end{aligned}$$

Both simulation of points and updates of their weights are trivial in this case. For our initial point

distribution I will take a bivariate uniform, to represent ignorance as to the location of x_0 . Plotted below is the uniform initial condition (green points) and the first set of generated points (blue points). The green points have uniform weight and the blue are sized in proportion to their new weight.




Figure 3: First Iteration of Example 2 SIS Filter

After 30 iterations I find that the filter is tracking the observations reasonably well, however I observe that $w_{30} = 0$ at all points except one (point in black in plot below). The green points are the initial condition, the red points are the observed path, and the blue points (plotted with uniform size which in this case is **not** representative of their weight) are the final points. The black point, which has $w = 1$ is totally *feasible* but the rate at which it accumulated all weight is suprising. However in the example in Elfring Et Al(2021) they describe 50% of points having zero weight after one iteration, and after five iterations all weights are zero except for one point. This is **very** similar to what I observe, and if weights decrease in this normal model like ours the decrease will be $\sim \exp(-(y-x)^2)$. So I think I'll just assume for now that what I've done is correct. I will move on and try and learn some more basics.

Final Thoughts For Now

At this point we have looked at a basic SIS Filter and written code for a simple example. I think I get the idea and I want to move on. We have also come across a couple of the issues described nicely by Elfring Et Al (2021). I will just list them here:

- **Degeneracy** is the term given to the phenomenon found above where a tiny number of points accumulate all weight. This leads to the point-weights providing very little information about the desired posterior, and estimated quantities will have high variance. Also, a huge fraction of computational work (recalcing weights) is wasted, as they are basically zero anyways. A solution is *resampling*, which we will look at later. An SIS filter with resampling




Figure 4: Thirty Iterations of Example 2 SIS Filter

are referred to as SISR.

- **Impoverishment** is the flipside to standard fixes for degeneracy. In resampling, we duplicate (with appropriate state noise) points with high weight and remove low weight points. This prevents degeneracy. However, if the state model variance is small we will end up with large numbers of essentially the same point. The results will be similar to the problem of degeneracy.
- **Divergence** is a catch-all expression for a range of problems resulting in the filter's inability to track the system generating the observations. See ElfringEtAl(2021) for a discussion of this point.
- **Choosing g_0** remains an obvious problem. The choice made above $g = p(x_t|x_{t-1})$ is a good one, but it is probably not possible in many cases. In that case, what to do?

We have covered the conceptual ideas required for particle filters, and more advanced topics appear to

- *practical work* involves fixing some of the problems we have already come across
- *theoretical work* involves investigating distributional properties, convergence, etc etc

Bayesian Filters for Static Parameters in Dynamic Models

The standard setup for filtering problems is the state-observation dynamics considered previously. An alternative framework considers the case of estimating static (ie time-invariant) parameters in a dynamically observed system. For example, assume we have an observation sequence $y_{1:T} = \{y_1, y_2, \dots, y_T\}$ from the model

$$dy = \kappa(y - \mu) + \sigma dW, \quad (24)$$

and we wish to compute online estimates of the parameters $\theta = (\kappa, \mu, \sigma)$. The relevant posterior to use here is $p(\theta|y_{1:t})$.

To work towards this we write $p(\theta|y_{1:t})$ recursively as done previously, deriving an expression $p(\theta|y_{1:t}) = k_{t-1,t}p(\theta|y_{1:t-1})$ as follows.

$$\begin{aligned} p(\theta|y_{1:t}) &= \frac{p(y_{1:t}|\theta) p(\theta)}{p(y_{1:t})} \\ &= \frac{p(y_t|y_{1:t-1}, \theta) p(y_{1:t-1}|\theta) p(\theta)}{p(y_t|y_{1:t-1}) p(y_{1:t-1})} \\ &= \frac{p(y_t|y_{1:t-1}, \theta) p(y_{1:t-1}|\theta) p(\theta)}{p(y_t|y_{1:t-1}) p(y_{1:t-1})} \\ &= k_{t-1,t} p(\theta|y_{1:t-1}) \end{aligned}$$

where

$$\begin{aligned} k_{t-1,t} &= \frac{p(y_t|y_{1:t-1}, \theta)}{p(y_t|y_{1:t-1})} \\ &= \frac{p(y_t|y_{1:t-1}, \theta)}{\int p(y_t|y_{1:t-1}, \theta) p(\theta|y_{1:t-1}) d\theta} \end{aligned}$$

is the required *transition kernel*. In this case the computational strategy is clearer after substituting $k_{t-1,t}$ back into the original expression to get

$$\begin{aligned} p(\theta|y_{1:t}) &= k_{t-1,t} p(\theta|y_{1:t-1}) \\ &= \frac{p(y_t|y_{1:t-1}, \theta) p(\theta|y_{1:t-1})}{\int p(y_t|y_{1:t-1}, \theta) p(\theta|y_{1:t-1}) d\theta} \end{aligned}$$

To solve using a particle approach we again approximate the posterior $p(\theta|y_{1:t})$ as a sequence of N points θ_t^i and weights $\pi_t^i = p(\theta^i|y_{1:t})$. To get there insert the particle approximation into the expression for $p(\theta|y_{1:t})$ above, and simplify using the Markov property assumed in our model

$$\begin{aligned} p(\theta|y_{1:t}) &= \frac{p(y_t|y_{1:t-1}, \theta) p(\theta|y_{1:t-1})}{\int p(y_t|y_{1:t-1}, \theta) p(\theta|y_{1:t-1}) d\theta} \\ &= \frac{p(y_t|y_{t-1}, \theta) p(\theta|y_{t-1})}{\int p(y_t|y_{t-1}, \theta) p(\theta|y_{t-1}) d\theta} \end{aligned}$$

and insert the particle approximation

$$\begin{aligned} p(\theta^i|y_t) &\sim \frac{p(y_t|y_{t-1}, \theta^i) p(\theta^i|y_{t-1})}{\sum p(y_t|y_{t-1}, \theta^i) p(\theta^i|y_{t-1})} \\ &= \frac{p(y_t|y_{t-1}, \theta^i) \pi_{t-1}^i}{\sum p(y_t|y_{t-1}, \theta^i) \pi_{t-1}^i} \\ &= \frac{\hat{\pi}_{t-1}^i}{\sum \hat{\pi}_{t-1}^i} \end{aligned}$$

where

$$\hat{\pi}_t^i = \pi_{t-1}^i p(y_t | y_{t-1}, \theta^i) \quad (25)$$

The process here is almost identical then to our previous examples. After initialising a set of N particles $\{\theta^i, \pi^i\}_{i=1}^N$ at time $t = 0$, we update our weights by computing the next $\hat{\pi}^i$ for each i using the above formula. The weights are then normalised to obtain the estimate of the posterior. The only difference here is that the particle points θ^i are not passed through a dynamic system like our model for states previously (because they are time-invariant). However, it may be the case that we modify them over time, but that is only a computational strategy and does not reflect any dynamics inherent to the model.